

# Training decision trees

Terence Parr  
MSDS program  
**University of San Francisco**

# Training overview

- Training partitions feature space into rectangular hypervolumes chasing reduced y impurity in subregions
- Hypervolumes are specified by sequence of *splits* that test a single feature and value at a time
- Each split becomes a decision node in decision tree
- Records in an “atomic” hypervolume form a single leaf
- Hypervolume described by conditionals on path from root to leaf

# How to create a decision node

- Each split chosen greedily to minimize impurity in subregion y's
  - Regressor: variance or MSE
  - Classifier: gini criterion or entropy
- To choose split, exhaustively try each (variable,value) pair and pick the pair with min weighted average impurity for subregions created by that split

# Fitting decision trees

subsets



MSE or gini



**Algorithm:** *dtreefit*(*X*, *y*, *min\_samples\_leaf*, *loss*)

**if**  $|X| < \text{min\_samples\_leaf}$  **then** return Leaf(*y*)

*col*, *split* = *bestsplit*(*X*, *y*, *loss*)

**if** *col* = -1 **then** return Leaf(*y*)    (*No better split?*)

*lchild* = *dtreefit*(*X*[*X*<sub>*col*</sub> ≤ *split*], *y*[*X*<sub>*col*</sub> ≤ *split*], *min\_samples\_leaf*, *loss*)

*rchild* = *dtreefit*(*X*[*X*<sub>*col*</sub> > *split*], *y*[*X*<sub>*col*</sub> > *split*], *min\_samples\_leaf*, *loss*)

**return** *DecisionNode*(*col*, *split*, *lchild*, *rchild*)

Overall fit: pass in full *X*, *y* to *dtreefit*() and get back the decision tree

# Best split var/value

**Algorithm:** *bestsplit*(*X*, *y*, *loss*)

*best* = (*col* = -1, *split* = -1, *loss* = *loss*(*y*))

**for** *col* = 1..*p* **do**

**foreach** *split* ∈ *X*<sub>*col*</sub> **do**

$y_l = y[X_{col} \leq split]$

$y_r = y[X_{col} > split]$

**if**  $|y_l| = 0$  **or**  $|y_r| = 0$  **then continue**

$l = \frac{|y_l| \times loss(y_l) + |y_r| \times loss(y_r)}{|y|}$       (weighted average of subregion losses)

**if**  $l = 0$  **then return** *col*, *split*

**if**  $l < best[loss]$  **then** *best* = (*col*, *split*, *l*)

**end**

**end**

**return** *best*[*col*], *best*[*split*]


Should pick midpoint between  
split value and next smallest x

# Decision tree prediction

```
1 Algorithm: predict(node, x)  
2   if node is leaf then  
3     if classifier then return mode(node.y)  
4     return mean(node.y)  
5   end  
6   if  $x[\text{node.col}] \leq \text{node.split}$  then return predict(node.lchild, x)  
7   return predict(node.rchild, x)
```


# The usual bestsplit() is inefficient

- It has a nested loop; tries all combinations of  $p$  variables and worst-case  $n$  unique values in each variable at root:  $O(n \cdot p)$
- Cost of computing loss on all values in subregion each iteration is also expensive
- For classification, can mitigate by sorting by  $i$ th var then we know at a specific  $x$  value, everything to left is less and right is greater; keep track of class counts to left/right
- Reduce computation by focusing on transitions points in  $x$ , effectively focusing on  $\text{unique}(x)$



# Improving generality and efficiency

- Select a subset of values as candidates,  $k$ ; then we reduce  $O(n \cdot p)$  to  $O(k \cdot p)$  for  $k \ll n$  ( $n$  is often huge) (our project  $k=11$ )
- We should really pick split point between two  $x$  values:  $(x^{(i)} + x^{(i-1)})/2$  (if sorted)
- More likely split point is between, not on,  $x$  values, so midpoint is good guess as to underlying distribution
- And, of course, we can reduce tree height with `min_samples_leaf` to restrict complexity





# Decision tree prediction via x subset

**Algorithm:** *bestsplit*( $X, y, loss$ )

$best = (col = -1, split = -1, loss = loss(y))$

**for**  $col = 1..p$  **do**

$candidates =$  randomly pick  $k \ll n$  values from  $X_{col}$

**foreach**  $split \in candidates$  **do**

$yl = y[X \leq split]$

$yr = y[X > split]$

**if**  $|yl| = 0$  **or**  $|yr| = 0$  **then continue**

$l = \frac{|yl| \times loss(yl) + |yr| \times loss(yr)}{|y|}$  (weighted average of subregion losses)

**if**  $l = 0$  **then return**  $col, split$

**if**  $l < best[loss]$  **then**  $best = (col, split, l)$

**end**

**end**

**return**  $best[col], best[split]$

Can even pick just 1 split randomly or in min..max range (see “Extremely random trees”); any small  $k$  value works.

# Prediction

- Start at the root node and descend through the decision nodes to the appropriate leaf
- At each decision node, test a specific variable's value against the split value stored in the decision node

```
1 Algorithm: predict(node, x)
2   if node is leaf then
3     if classifier then return mode(node.y)
4     return mean(node.y)
5   end
6   if  $x[\text{node.col}] \leq \text{node.split}$  then return predict(node.lchild, x)
7   return predict(node.rchild, x)
```